

# Measuring Software Ticket Quality using Quantitative Data Analysis

Andrei-Mihai Nicolae (2147392)  
Dr. Timothy Storer

# Introduction

- Software engineering is becoming more complex as technology has become ubiquitous
- Harder to plan, manage and track work during development lifecycle [8]
- Solution? **Issue tracking systems** (e.g. Jira, Bugzilla, Manuscript)
- **Software Tickets** are the core component of such systems

# Software Tickets

- Many components: summary, description, attachments, comments...
- Usually come in one of 2 forms: feature requests and bug reports
- What makes for a **High Quality** software ticket?
- Our findings - contribution to the community

The screenshot shows a software ticket interface for 'Angry Nerds / ANGRY-304' titled 'Red Angry Nerd is scary'. The interface includes a top navigation bar with buttons for 'Edit', 'Comment', 'Assign', 'More', 'Start Progress', 'Resolve Issue', 'Workflow', and 'Admin'. Callouts highlight the 'Log work, attach files & screenshots, create sub-tasks, move, link, or clone the issue.' (pointing to the 'More' button), 'Transition the issue in its workflow here' (pointing to the 'Workflow' button), 'Add a field or access the Admin helper from this menu' (pointing to the 'Admin' button), and 'Email this issue to others' (pointing to the 'Email' button). The 'Details' section shows the ticket type as 'Bug', priority as 'Low', and status as 'Waiting for Triage'. The 'Description' section has a placeholder 'Click to add description'. The 'Attachments' section shows a file named 'hydra.jpg' (67 kB, 21/Mar/13 3:38 PM). The 'People' section lists assignee 'Susan Griffin', reporter 'Bartek Gatz', and options to 'Vote for this issue' and 'Start watching this issue'. The 'Dates' section shows 'Created: 21/Mar/13 3:37 PM', 'Updated: 16/May/13 3:36 PM', 'Scheduled: 21/Mar/13', and 'Deployment Date:'. The 'Development' section has a 'Create Branch' button, with a callout 'To create a code branch in Bitbucket or Stash, click the Create Branch link'. The 'Agile' section has a 'View on Board' link. The 'Analytics' section is empty. The 'Activity' section at the bottom has tabs for 'All', 'Comments', 'History', 'Activity', 'Source', 'Reviews', 'Transitions Summary', 'Commits', and 'Builds'.

Angry Nerds / ANGRY-304  
Red Angry Nerd is scary

Log work, attach files & screenshots, create sub-tasks, move, link, or clone the issue.

Transition the issue in its workflow here

Add a field or access the Admin helper from this menu

Email this issue to others

Export this issue to other formats, such as MS Word

1 of 8  
Return to Search

Edit Comment Assign More Start Progress Resolve Issue Workflow Admin

Details

Type: Bug  
Priority: Low  
Component/s: None  
Labels: None  
Monkey: Cheeky Monkey

Status: Waiting for Triage (View Workflow)  
Resolution: Unresolved  
Fix Version/s: None

People

Assignee: Susan Griffin  
Reporter: Bartek Gatz  
Votes: 0 Vote for this issue  
Watchers: 1 Start watching this issue

Description

Click to add description

Attachments

hydra.jpg  
67 kB 21/Mar/13 3:38 PM

Dates

Created: 21/Mar/13 3:37 PM  
Updated: 16/May/13 3:36 PM  
Scheduled: 21/Mar/13  
Deployment Date:

Development

Create Branch

To create a code branch in Bitbucket or Stash, click the Create Branch link

Agile

View on Board

Analytics

Activity

All Comments History Activity Source Reviews Transitions Summary Commits Builds

# Contributions

1. Innovative Go tool built for providing efficient data collection and analysis; open sourced on GitHub [2]
2. One of the few studies in the field that performs a quantitative analysis rather than a qualitative one
3. One of the very few research projects that investigates such a large number of tickets (over 300,000) extracted from 38 different projects
4. To our knowledge, the first study to conduct sentiment and grammar correctness analyses on software tickets

# Research Questions

- Does the presence of attachments and their type (e.g. code snippet, screenshot) influence the **Time-To-Close** for a ticket?
  - Does the presence of stack traces reduce **Time-To-Close**?
  - Does the presence of steps to reproduce reduce **Time-To-Close**?
  - Is there a relationship between the number of words in comments and **Time-To-Close**?
- Does the total number of words in summary and description have an impact on **Time-To-Close**?
  - Does the number of grammar errors in summary, description and comments have an effect on **Time-To- Close**?
  - Does a positive or negative ticket influence its **Time- To-Close**?

# Related Work

- Bettenburg et al. [1] - qualitative analysis through interviewing developers about what makes for high quality tickets; developed Cuezilla for predicting quality
- Hooimeijer et al. [3] - analyzed over 25,000 tickets; found that readability, attachments and comments can significantly increase the quality of a ticket
- Schroeter et al. [4] - investigated the effect stack traces have on ticket lifespan; around 60% of tickets with stack traces were fixed in one of the methods in the frame, 40% in the first frame
- Bettenburg et al. [5] and Prifti et al. [6] looked at bug report duplicates and their consequences; found that they actually bring value to the project, duplicates usually offering extra information not found in master report

# Building the Data Set

- Needed a tool for fetching, storing, analyzing, plotting graphs and running statistical tests on tickets
- Should be fast
- Should have support for multiple databases
- Should have great HTTP support
- Should have plenty of libraries
- Solution? Go application called Ticket Guru (amazing mascot to the right 🐿)

The Almighty Ticket Guru

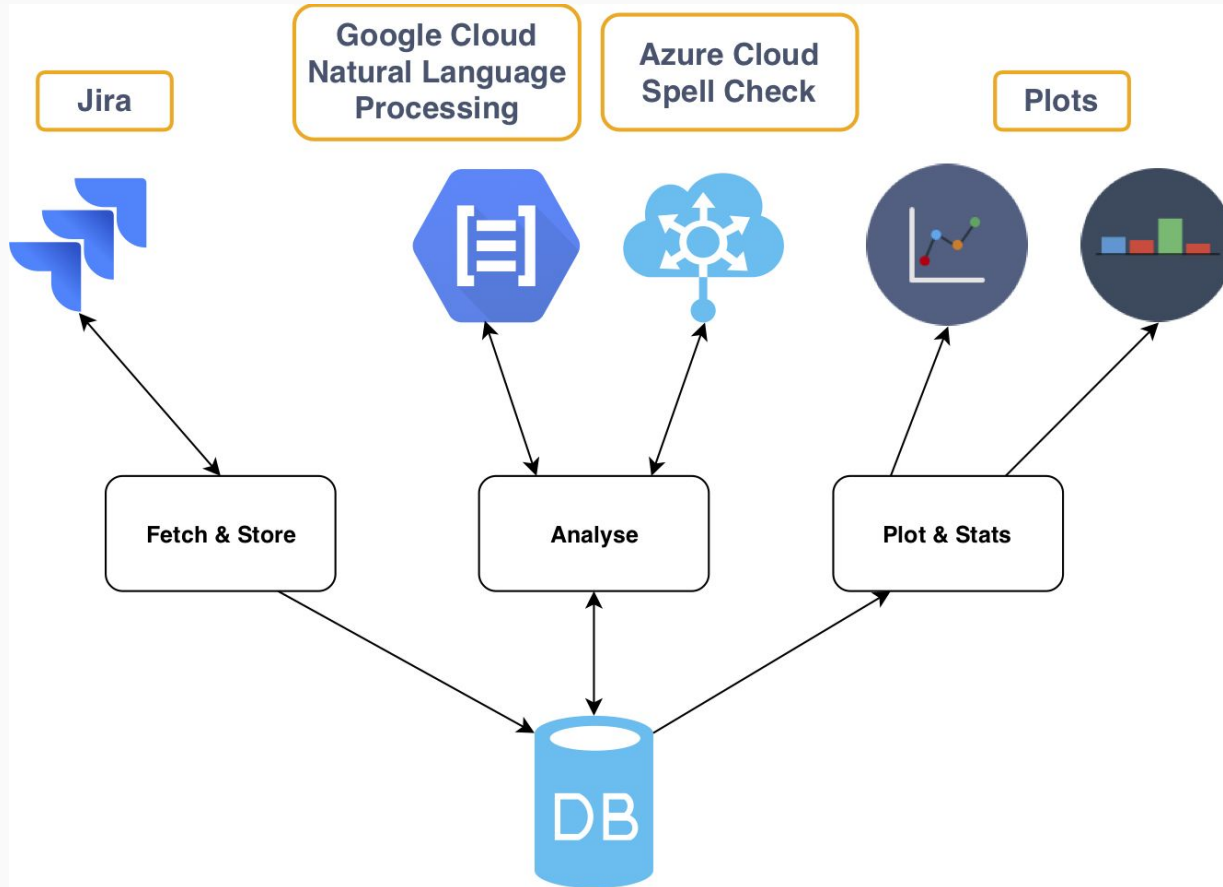


# Ticket Guru Structure





# Ticket Guru Flow



# Characterizing the Data Set

**303,138**

Total number of tickets

- All collected from Apache Jira
- 38 different projects
- Different programming languages, ranging from Java to Python and Ruby

**236,383**

Closed tickets

Tickets marked **Closed**, **Resolved**, **Done** or **Completed** when they were fetched.

**201,786**

Tickets eligible for analysis

- Closed tickets
- High Priority (i.e. Critical, Blocker, Major, High)
- No outliers

# Characterizing the Data Set

**270,907**

**Tickets with comments**

With at least one comment in the discussion thread of the ticket.

**39,988**

**Tickets with Steps-To-Reproduce**

Used complex regex [7] to determine whether the tickets had Steps-To-Reproduce in summary, description or comments.

**1,942**

**Tickets with Java stack traces**

Made use of technique outlined by Bettenburg et. al [7] to determine whether tickets had stack trace(s) in description or comments.

# Characterizing the Data Set

**157,047**

**Tickets with sentiment  
scores**

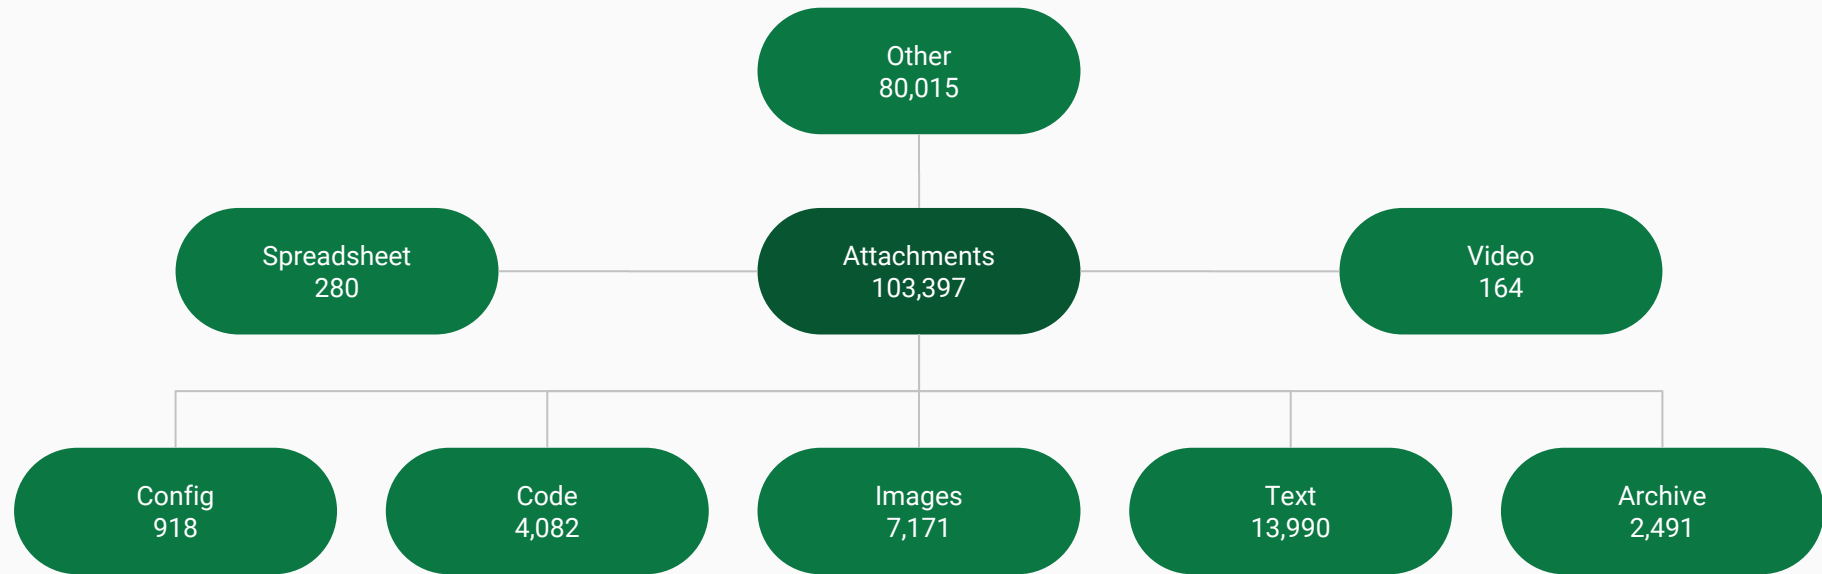
Sentiment scores were calculated using  
Google's Natural Language Processing API.

**133,689**

**Tickets with grammar  
correctness scores**


Grammar correctness scores were  
calculated using Azure Cloud Bing Spell  
Check API.

# Characterizing the Data Set




# Correlations

## Attachments analysis




## Comments Complexity Analysis




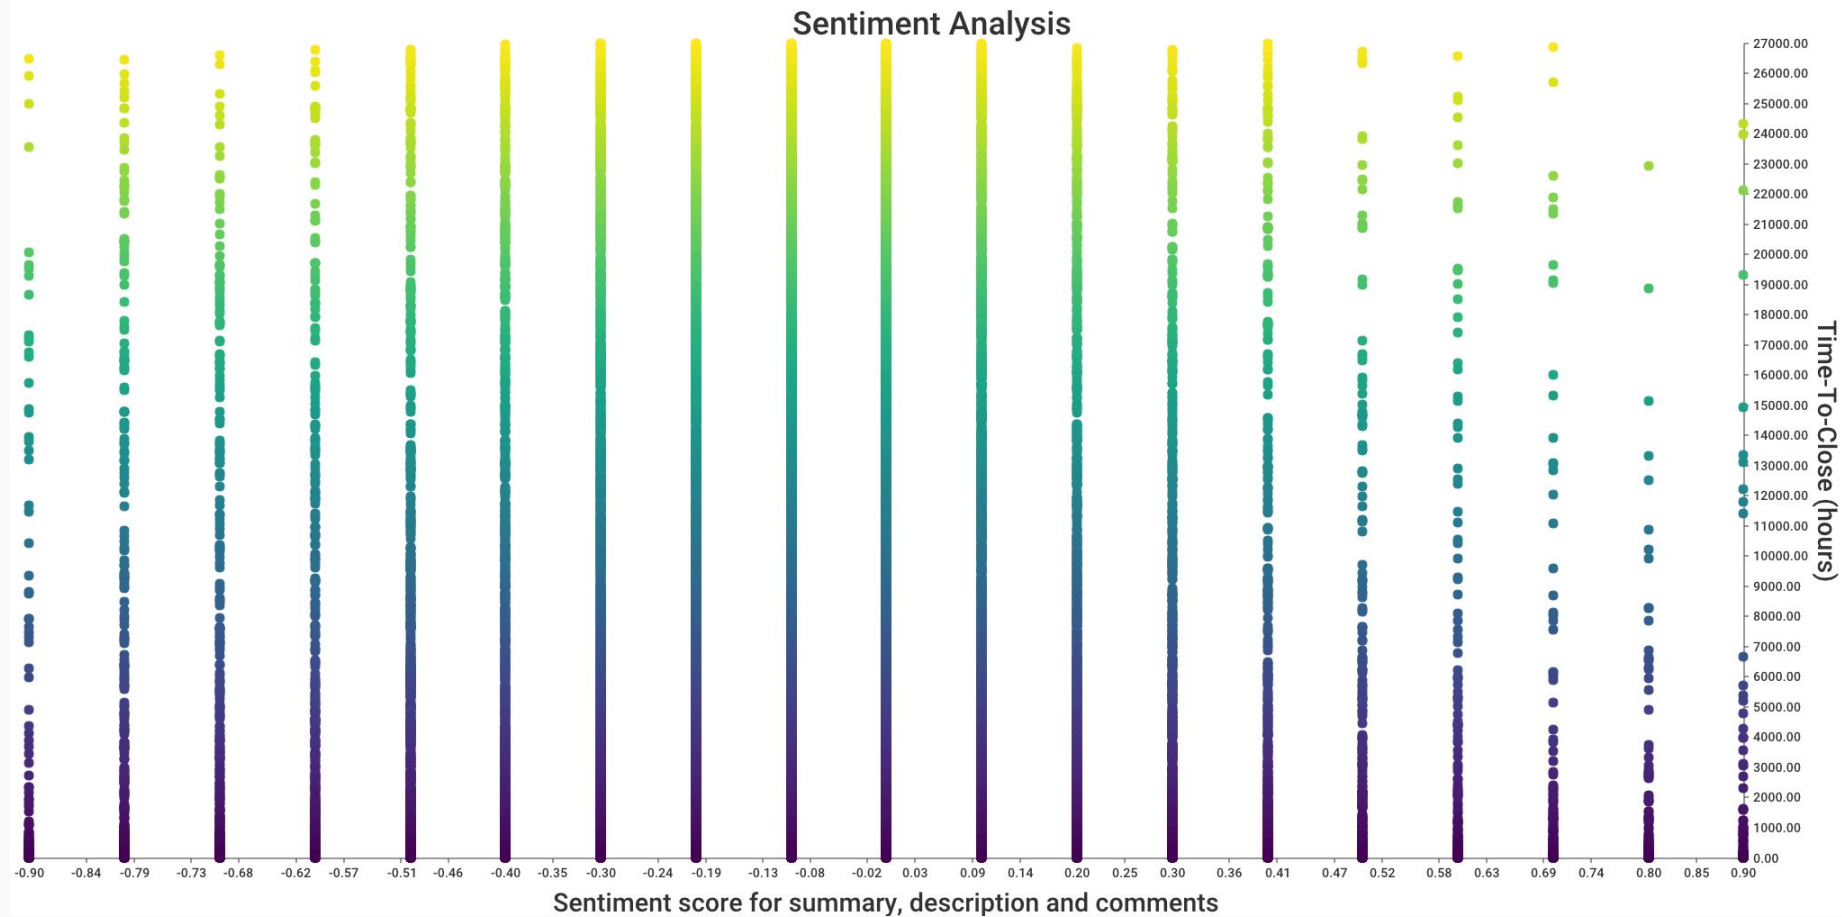


## Fields Complexity Analysis




## Grammar Correctness Analysis






## Stack Traces Analysis



## Steps To Reproduce Analysis



# Future Work

- Obtain Google Cloud Platform credits for more comprehensive analysis
- Determine ticket difficulty from information inside it - currently not known how to do it
- Create a *goodness* metric or a *recommender tool* that can automatically create high quality tickets
- Add support for Bugzilla
- Increase the current database of tickets
- Include closed source projects in the analysis

# Conclusions

- Software Tickets are a vital part of every software project lifecycle
- Deriving quality score is not trivial
- Answered all RQs with statistically significant results
- Bring valuable contributions to future work in this area

Thank you!





1. N. Bettenburg, S. Just, A. Schroeter, C. Weiss, R. Premraj, and T. Zimmermann. What makes a good bug report? Pages 308–318, 2008
2. <https://github.com/nclandreiticketguru>
3. P. Hooimeijer and W. Weimer. Modeling bug report quality. pages 34–43, 2007
4. N. Bettenburg, S. Just, A. Schroeter, C. Weiss, R. Premraj, and T. Zimmermann. What makes a good bug report? pages 308–318, 2008
5. N. Bettenburg, R. Premraj, T. Zimmermann, and S. Kim. Duplicate bug reports considered harmful: really? pages 337–345, 2008
6. T. Prifti, S. Banerjee, and B. Cukic. Detecting bug duplicate reports through local references. In Proceedings of the 7th International Conference on Predictive Models in Software Engineering, page 8. ACM, 2011
7. N. Bettenburg, R. Premraj, T. Zimmermann, and S. Kim. Extracting structural information from bug reports. pages 27–30, 2008
8. J. D. Herbsleb. Global software engineering: The future of socio-technical coordination. In 2007 Future of Software Engineering, pages 188–198. IEEE Computer Society, 2007