

Measuring Software Ticket Quality using Quantitative Data Analysis

Andrei-Mihai Nicolae (2147392)
Dr. Timothy Storer

Introduction

- Software engineering is becoming more complex as technology has become ubiquitous
- Harder to plan, manage and track work during development lifecycle [8]
- Solution? **Issue tracking systems** (e.g. Jira, Bugzilla, Manuscript)
- **Software Tickets** are the core component of such systems

Software Tickets

- Many components: summary, description, attachments, comments...
- Usually come in one of 2 forms: feature requests and bug reports
- What makes for a **High Quality** software ticket?
- Our findings = contribution to the community

The screenshot shows a software ticket interface for 'Red Angry Nerd is scary' (ANGRY-304). The interface includes a top navigation bar with buttons for 'Edit', 'Comment', 'Assign', 'More', 'Start Progress', 'Resolve Issue', 'Workflow', and 'Admin'. Callouts highlight the 'Log work, attach files & screenshots, create sub-tasks, move, link, or clone the issue.' and 'Transition the issue in its workflow here'.

The ticket details section shows the following information:

- Type: Bug
- Priority: Low
- Component/s: None
- Labels: None
- Monkey: Cheeky Monkey
- Status: Waiting for Triage (View Workflow)
- Resolution: Unresolved
- Fix Version/s: None

The description section is empty, with a link to 'Click to add description'.

The attachments section shows a file named 'hydra.jpg' (67 kB) uploaded on 21/Mar/13 3:38 PM.

The right sidebar contains sections for 'People' (Assignee: Susan Griffin, Reporter: Bartek Gatz), 'Dates' (Created: 21/Mar/13 3:37 PM, Updated: 16/May/13 3:36 PM, Scheduled: 21/Mar/13, Deployment Date: 21/Mar/13), 'Development' (Create Branch), 'Agile' (View on Board), and 'Analytics'.

Callouts highlight the 'Email this issue to others' and 'Export this issue to other formats, such as MS Word'.

The bottom section shows the 'Activity' tab with sub-tabs for 'All', 'Comments', 'History', 'Activity', 'Source', 'Reviews', 'Transitions Summary', 'Commits', and 'Builds'.

Contributions

1. Innovative Go tool built for providing efficient data collection and analysis; open sourced on GitHub [2]
2. One of the few studies in the field that performs a quantitative analysis rather than a qualitative one
3. One of the very few research projects that investigates such a large number of tickets (over 300,000) extracted from 38 different projects
4. To our knowledge, the first study to conduct sentiment and grammar correctness analyses on software tickets

Research Questions

- Does the presence of attachments and their type (e.g. code snippet, screenshot) influence the **Time-To-Close** for a ticket?
- Does the presence of stack traces reduce **Time-To-Close**?
- Does the presence of steps to reproduce reduce **Time-To-Close**?
- Is there a relationship between the number of words in comments and **Time-To-Close**?
- Does the total number of words in summary and description have an impact on **Time-To-Close**?
- Does the number of grammar errors in summary, description and comments have an effect on **Time-To-Close**?
- Does a positive or negative ticket influence its **Time-To-Close**?

Related Work

- Bettenburg et al. [1] - qualitative analysis through interviewing developers about what makes for high quality tickets; developed Cuezilla for predicting quality
- Hooimeijer et al. [3] - analyzed over 25,000 tickets; found that readability, attachments and comments can significantly increase the quality of a ticket
- Schroeter et al. [4] - investigated the effect stack traces have on ticket lifespan; around 60% of tickets with stack traces were fixed in one of the methods in the frame, 40% in the first frame
- Bettenburg et al. [5] and Prifti et al. [6] looked at bug report duplicates and their consequences; found that they actually bring value to the project, duplicates usually offering extra information not found in master report

Building the Data Set

- Needed a tool for fetching, storing, analyzing, plotting graphs and running statistical tests on tickets
- Should be fast
- Should have support for multiple databases
- Should have great HTTP support
- Should have plenty of libraries
- Solution? Go application called Ticket Guru (amazing mascot to the right 🐻)

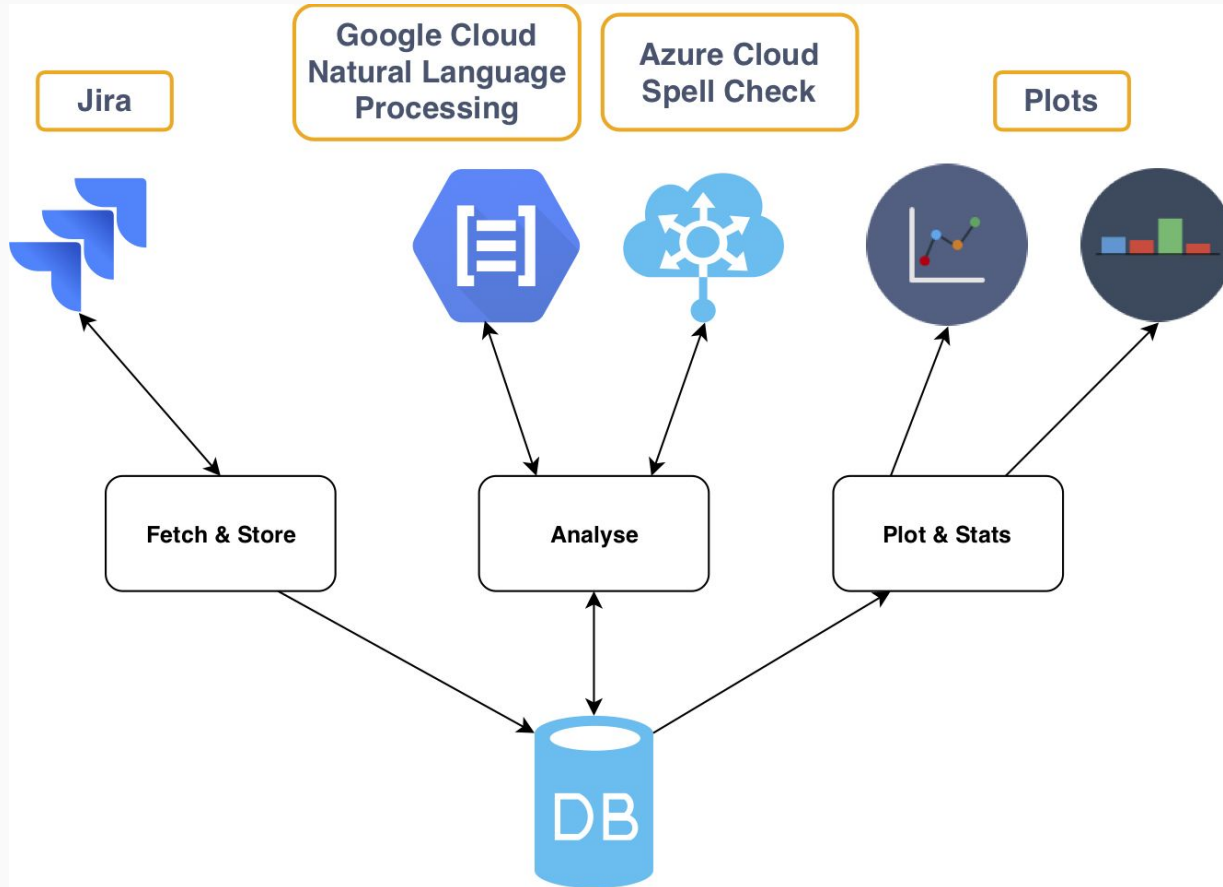
The Almighty Ticket Guru



Ticket Guru Structure



Ticket Guru Flow



Characterising the Data Set (1)

303,138

Total number of tickets

- All collected from Apache Jira
- 38 different projects
- Different programming languages, ranging from Java to Python and Ruby

236,383

Closed tickets

Tickets marked **Closed**, **Resolved**, **Done** or **Completed** when they were fetched.

201,786

Tickets eligible for analysis

- Closed tickets
- High Priority (i.e. Critical, Blocker, Major, High)
- No outliers

Characterising the Data Set (2)

113,344

Tickets with comments

With at least one comment in the discussion thread of the ticket.

39,988

Tickets with Steps-To-Reproduce

Used complex regex [7] to determine whether the tickets had Steps-To-Reproduce in summary, description or comments.

1,942

Tickets with Java stack traces

Made use of technique outlined by Bettenburg et. al [7] to determine whether tickets had stack trace(s) in description or comments.

Characterising the Data Set (3)

157,047

Tickets with sentiment scores

Sentiment scores were calculated using Google's Natural Language Processing API.

133,689

Tickets with grammar correctness scores

Grammar correctness scores were calculated using Azure Cloud Bing Spell Check API.


98,311

Tickets with attachments


Tickets that had one of the following types of attachments: code, configuration, image, video, spreadsheet, text, archive or other.

Correlations


Attachments analysis




Comments Complexity Analysis

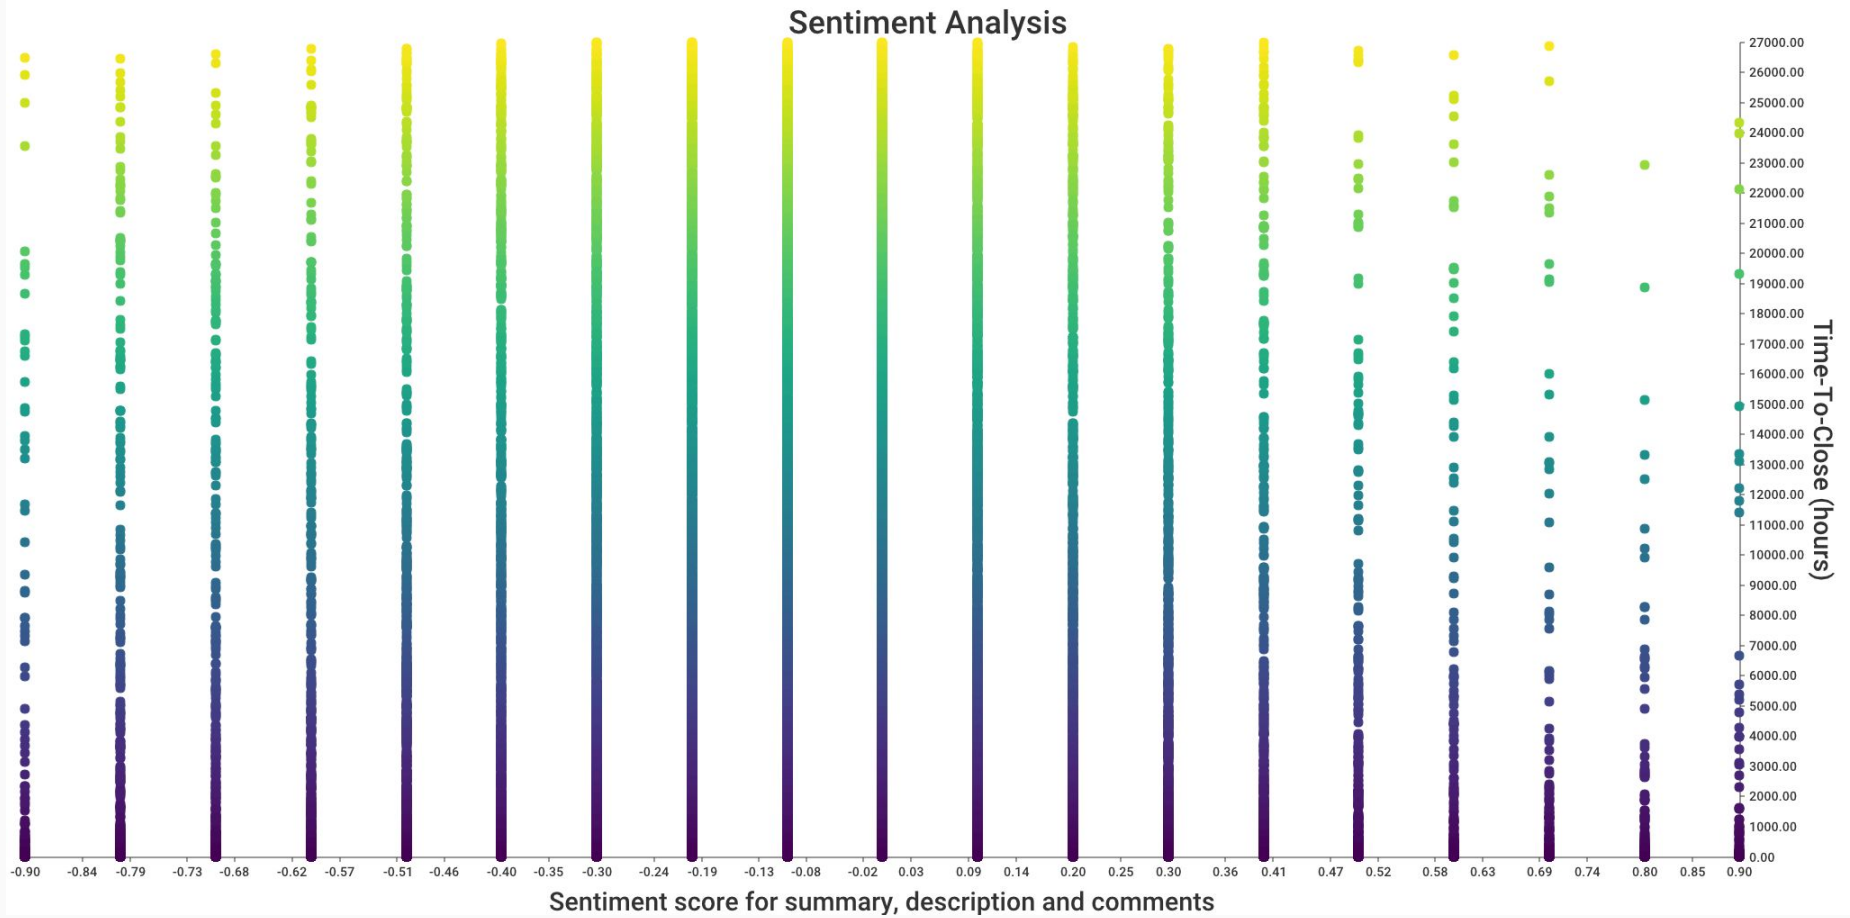


Fields Complexity Analysis




Grammar Correctness Analysis






Stack Traces Analysis



Steps To Reproduce Analysis



Future Work

- Obtain Google Cloud Platform credits for more comprehensive analysis
- Determine ticket difficulty from information inside it - currently not known how to do it
- Create a *goodness* metric or a *recommender tool* that can automatically create high quality tickets
- Add support for Bugzilla
- Increase the current database of tickets
- Include closed source projects in the analysis

Conclusions

- Software Tickets are a vital part of every software project lifecycle
- Deriving quality score is not trivial
- Answered all RQs with statistically significant results
- Bring valuable contributions to future work in this area

Thank you!



1. N. Bettenburg, S. Just, A. Schroeter, C. Weiss, R. Premraj, and T. Zimmermann. What makes a good bug report? Pages 308–318, 2008
2. <https://github.com/nclandreiticketguru>
3. P. Hooimeijer and W. Weimer. Modeling bug report quality. pages 34–43, 2007
4. N. Bettenburg, S. Just, A. Schroeter, C. Weiss, R. Premraj, and T. Zimmermann. What makes a good bug report? pages 308–318, 2008
5. N. Bettenburg, R. Premraj, T. Zimmermann, and S. Kim. Duplicate bug reports considered harmful: really? pages 337–345, 2008
6. T. Prifti, S. Banerjee, and B. Cukic. Detecting bug duplicate reports through local references. In Proceedings of the 7th International Conference on Predictive Models in Software Engineering, page 8. ACM, 2011
7. N. Bettenburg, R. Premraj, T. Zimmermann, and S. Kim. Extracting structural information from bug reports. pages 27–30, 2008
8. J. D. Herbsleb. Global software engineering: The future of socio-technical coordination. In 2007 Future of Software Engineering, pages 188–198. IEEE Computer Society, 2007