



University
of Glasgow | School of
Computing Science

Software Ticket Quality

Andrei-Mihai Nicolae

School of Computing Science
Sir Alwyn Williams Building
University of Glasgow
G12 8QQ

Masters project proposal

December 18th 2017

Contents

1	Introduction	2
1.1	A subsection	2
2	Statement of Problem	2
3	Background Survey	2
3.1	Introduction	2
3.2	What was covered and why	2
3.3	How papers were found; what terms were searched for + snowballing . . .	2
3.4	Data Quality Metrics	2
3.5	Issue Quality	8
3.6	Measuring Cost and Waste in Software Projects	20
3.7	Sentiment Analysis	22
3.8	Conclusion	23
4	Proposed Approach	23
5	Work Plan	23

1 Introduction

briefly explain the context of the project problem

1.1 A subsection

Please note your proposal need not follow the included section headings; this is only a suggested structure. Also add subsections etc as required

2 Statement of Problem

Clearly state the problem to be addressed in your forthcoming project. Explain why it would be worthwhile to solve this problem.

3 Background Survey

3.1 Introduction

3.2 What was covered and why

3.3 How papers were found; what terms were searched for + snowballing

3.4 Data Quality Metrics

Software Process Data Quality and Characteristics[5]:

- **What were the goals?** The authors are trying to investigate different software process data (e.g. bug tracking databases, version control systems logs) in order to infer quality and characteristics. Then, they want to compute measures and discuss the issues that arise from these observations.
- **What was the method?** They examined 6 software projects, 5 of which were open source, by extracting the SVN and CVS logs, as well as the content of the bug tracker databases. Then, they tried to link the logs with the contents of the bug tracker databases as they are not integrated by default.

- **What did they learn?** Firstly, the closed source project the authors examined showed much more data quality as the open source counterparts, among which better average attachments per bug report, better average status changes per bug report or average commits per bug report/developer. Moreover, the authors found that there is little linkage between bug reports and version control logs, thus reducing the overall traceability. Also, there is reduced quality in change logs as, for example, the Eclipse project has more than 20% empty commit messages. In conclusion, closed source projects will probably exhibit improved data quality than open source ones, but data quality in general differs from project to project.
- **Relevance to our work:** The analysis on data quality in open/closed source projects presented in this research is valuable to us as it presents various statistics on commits, bug report structure and quality, as well as linkage between the two. We might apply the techniques in our own research on ticket quality.

Data Quality in Context[34]:

- **What were the goals?** The main goal of the authors was to analyze data quality in different organizations, observe data quality problem patterns as well as see its implications.
- **What was the method?** They performed a qualitative analysis, examining 42 data quality projects from 3 organizations. They collected data via interviews of data procedures, custodians, customers and managers. The authors analysed each project using the data quality dimensions as content analysis codes.
- **What did they learn?** One thing that the authors learned is that representational data quality dimensions are underlying causes of accessibility data quality problem patterns. Also, they observed three underlying causes for users' complaints regarding data not supporting their tasks:
 - incomplete data;
 - inadequately defined or measured data;
 - data that could not be appropriately aggregated.
- **Relevance to our work:** this paper brings valuable insight into major data quality issues and patterns, information which we can apply in our own research.

Data Quality Assessment[27]:

- **What were the goals?** The main goal of the authors is to present ways of developing usable data quality metrics that organizations can implement in their own internal processes.

- **What was the method?** They have researched previous methods of assessing data quality in information systems and they formulated general advice that can be directly implemented by organizations in their processes.
- **What did they learn?** They learned that there is no universal approach to assessing data quality as it heavily depends on the context where it is analysed.
- **Relevance to our work:** it shows how one can effectively measure data quality in information systems, thus we can apply it on the quality of the tickets we are analysing.

Product perspective on total Data Quality Assessment[36]:

- **What were the goals?** The author wanted to create a Total Data Quality Management methodology so that high quality information products could be delivered to information consumers.
- **What was the method?** The research presents us, firstly, with a couple of definitions for the methodology:
 - TDQM cycle - define, measure, analyze and improve information quality continuously;
 - information product - information output from an information manufacturing system that has value to the information consumer;
 - information quality - fitness for use by information consumers.

Then, the author presents the four steps in order to successfully implement the TDQM methodology:

- clearly articulate the information product;
 - establish several roles that would be in charge of the information product management;
 - teach information quality assessment and management skills to all IP constituencies;
 - institutionalize continuous information product improvement.
- **What did they learn?** The author learned that this methodology is filling a gap in the information quality world and that it could help organizations better manage the information quality of their information products.
- **Relevance to our work** This article could prove helpful to our own research as the methodology presented could help us in detecting if there is data quality in the projects we will be analyzing.

Knowing-Why about Data Processes and Data Quality[23]:

- **What were the goals?** They tried to investigate whether knowing-why affects work performance and whether knowledge held by people in different roles affect the overall work performance.
- **What was the method?** They had 6 companies that served as data collection points. In all of them, the research focused on the quality of their customer activity data, tuned for the specific domain the company was activating in. They held 2 sessions at each company: in the first session, the participants were presented with an overview of the research and a questionnaire, while during the second session provided summary feedback on the questionnaire results and gave the participants the chance to provide comments and explanations.
- **What did they learn?** The major finding of the study was that there is vast complexity of knowledge at work in data production process.
- **Relevance to our work:** an important research on how knowledge in different roles in a company can lead to decreased data quality overall. It will help us to analyse the software tickets we will mine better.

Antecedents of Information and System Quality [26]:

- **What were the goals?** The authors had 2 main research objectives they wanted to aim for in this paper:
 - identify a set of antecedents that both drive information and system quality, as well as define the nature of the IT artifact;
 - explore data warehousing in general, especially analytical tools, predefined reports and ad hoc queries.
- **What was the method?** They conducted a cross-functional survey to test the model the authors proposed. They focused on user experience with a data warehouse. They sent email invites to the Data Warehouse Institute which distributed it further to organizations, among which 7 agreed to participate, yielding a total of 465 completed surveys.
- **What did they learn?** They learned that their selected features were indeed a good indicator of overall information and system quality. They also found out that accuracy is the dominant determinant across all three data warehouse technologies, followed by completeness and format. Also, the results showed that more attention needs to be given to differences across varying technologies.
- **Relevance to our work:** this paper presents important findings regarding information and system quality in an IT environment. We can use it when analysing the quality of the tickets we will mine.

AIMQ[24]:

- **What were the goals?** The main challenge set by the authors was to develop an overall model with an accompanying assessment instrument for quantifying information quality.
- **What was the method?** They developed a methodology called AIM quality (hence the name AIMQ) which they applied at 5 different organizations. It has 3 main components:
 - 2×2 model of what information quality means to managers;
 - questionnaire for measuring information quality along the dimensions found in first step;
 - two analysis techniques for interpreting the assessments captured by the questionnaire.
- **What did they learn?** They conclude that the AIMQ methodology proved to be a practical information quality tool to organizations.
- **Relevance to our work:** the paper gave us valuable insight into how to assess information quality.

Evaluating IT Quality in E-Government Environment[28]:

- **What were the goals?** The main goal of the paper is to identify if leadership and IT quality can lead to positive outcomes in an e-government environment.
- **What was the method?** They conducted a web survey to gather data and test their hypotheses. They assessed the City of Denton's e-government initiatives, including current plans and implementations.
- **What did they learn?** They learned that the MBNQA leadership triad (leadership, strategic planning and customer/market focus) had a positive impact on the IT quality triad (information, system and service quality). Moreover, they found out that both leadership and IT quality improved the benefits.
- **Relevance to our work:** even though the environment presented in the paper was an e-government, it is useful to see how the organization around a project as well as IT quality can lead to a better functioning of the system. Thus, in our research, it can prove helpful in choosing the right open source project to gather and analyse the data from.

Organizational Impact of Information Quality[12]:

- **What were the goals?** The authors had two main research questions they wanted to address:
 - what individual/combined influences system quality, service quality, information quality have on organizational impact?

- what effect does system quality have on information quality?
- **What was the method?** They conducted a construct measurement for the IS quality dimensions and collected data through empirical testing.
- **What did they learn?** Firstly, they learned that service quality has the greatest impact of all three quality constructs. Moreover, they found out that there is a linkage between system and information quality, fact that previous research did not reveal. Lastly, their results indicate that the above-mentioned quality dimensions have a significant positive influence on organizational impact either directly or indirectly.
- **Relevance to our work:** the paper can provide us insights into how valuable information quality is and what an impact it can have.

Impact of Poor Data Quality[30]:

- **What were the goals?** The main purpose of the article is to provide some insights into how data quality can have a major economical impact on a company, as well as to raise awareness about this issue.
- **What was the method?** The author has analysed various large enterprises and their available data to infer conclusions about the impact data quality has.
- **What did they learn?** There are three main issues across most enterprises:
 - inaccurate data;
 - inconsistencies across databases;
 - unavailable data necessary for certain operations.

Moreover, the paper also presents a list of impacts poor data quality brings, among which we can find:

- lowered customer satisfaction;
- increased cost;
- lowered employee satisfaction;
- it becomes more difficult to set strategy;
- poorer decision making.
- **Relevance to our work:** this is a key paper on data quality in an enterprise setting and what impact it can have. It is valuable to our research because we can relate the findings with what negative impacts poor tickets can bring to software projects.

Metrics for Measuring Data Quality[15]:

- **What were the goals?** The authors of the paper aimed to answer the two main research questions:

- how can data quality be measured as metrics?
- what can be done in regards to data quality and what economic consequences would it bring?
- **What was the method?** They applied the metric for timeliness at a major German mobile services provider. Due to some Data Quality issues the company was having, they had lower mailing campaign success rates, but after applying the metrics, the company was able to establish a direct connection between the results of measuring data quality and the success rates of campaigns.
- **What did they learn?** For some practical problems, the metrics that they applied in regards to data quality proved to be quite appropriate. As they designed the metrics to revolve around interpretability and cardinality, they could quantify data quality, thus they could analyse the economical impact.
- **Relevance to our work:** this paper could possibly help us in measuring data quality. This is an important step in our research as we specifically need to measure quality in software tickets.

3.5 Issue Quality

Using fuzzy code search to link code fragments in discussions to source code[10]:

- **What were the goals?** The main goal of the paper is to find an approach that improves on previous methods to link source code fragments in discussions to source code. These previous methods were mainly based on information retrieval or change log analysis.
- **What was the method?**
- **What did they learn?**
- **Relevance to our work:**

Where should we fix this bug?[20]:

- **What were the goals?** The main goal of the paper was to aid developers in easily finding where to apply the necessary changes in order to fix a bug. The authors wanted to bridge this gap by creating a two-phase recommendation model that would locate the locations where the fixes need to be applied based on the corresponding bug reports.
- **What was the method?** They divided the work into multiple stages:
 - firstly, they did a feature extraction on the bug reports (e.g. extract description, metadata, summary);

- in order to predict the locations of the necessary fixes, the one-phase recommendation model was put into place - this model would be trained on collected bug reports and then, when given a new bug report, it would try to localize the files that need to be changed automatically;
- last step was to implement the two-phase recommendation model which basically has two classification phases - binary and multiclass:
 - * binary phase filters out uninformative bug reports before predicting the files to fix;
 - * multiclass phase where the previously filtered bug reports are used as training data and only after that new bug reports can be analyzed and files to be changed recommended.
- **What did they learn?** Evaluating the model yielded a 70% accuracy in pointing to the correct files to be fixed, as long as the algorithm actually recommended any locations. However, the authors acknowledge that the approach needs extra work, especially in fine-grained defect localization, performance improvement and bug triaging.
- **Relevance to our work:** The research shows interesting points regarding time-to-fix-bug and bug report contents, thus it can give us an insight into how ticket quality can affect that time to fix the bug (in the case of this paper it would mean to successfully pass the first phase of the recommendation model so that fix locations could be proposed).

Do Stack Traces Help Developers Fix Bugs?[31]:

- **What were the goals?** The main goal of the paper was to identify if including a stack trace in a bug report would help the developers fix the issue quicker. They divided their investigations in order to follow 4 research questions:
 - are bugs fixed in methods in stack traces;
 - how far down the stack to trace;
 - are two or more stack traces better than one;
 - do stack traces help speed up debugging.
- **What was the method?** They conducted their investigation on the Eclipse project which is developed in Java. Then, they extracted stack traces from bug reports using the infoZilla tool. Afterwards, they linked the stack traces (and thus the corresponding bug report) to changes in the source code (change logs) by mining the version repository of the Eclipse project.
- **What did they learn?** They found answers to all their research questions, and are as follows:
 - around 60% of the bugs that contained stack traces were fixed in one of the methods in the frames;

- around 40% were fixed in the first stack frame, around 80% were fixed within the top 6 frames, while close to 90% of the bugs were fixed within top 10 stack frames;
 - having more than one stack trace usually helped more than only a single one;
 - bug reports containing stack traces usually have a shorter lifetime, thus speeding up the debugging process.
- **Relevance to our work:** Stack traces are one of the types of information we are investigating in our research and it is very useful to have such a study back up the statement that stack traces are valuable for developers so that they speed up the issue fixing process.

Extracting Structural Information from Bug Reports[9]:

- **What were the goals?** The authors try to present a tool called infoZilla that can accurately extract different types of information (e.g. stack traces, source code fragments) from bug reports.
- **What was the method?** They selected a few categories of information that both infoZilla could successfully extract from bug reports and that are of most interest to developers. Thus, the categories selected are:
 - patches;
 - stack traces;
 - source code;
 - enumerations.

Then, they ran the tool on tickets and evaluated the overall accuracy.

- **What did they learn?** After developing the tool to inspect such elements and retrieve them, they evaluated it on over 160000 Eclipse bug reports. For all categories, infoZilla achieved over 97% accuracy.
- **Relevance to our work:** If we could manage to get this tool and use it in our own research, it would be beneficial as we could extract efficiently different types of information we are interested in testing ticket quality (e.g. stack traces).

Automatic Bug Report Assignment[1]:

- **What were the goals?** The authors main goal was to determine a way through which bug triaging can be made easier using automatization approaches.

- **What was the method?** The researchs output was an algorithm that can automatically assign bugs to developers based on various types of information fed. The author looked exclusively at ML approaches, examining 8 main types of information (textual description of bug, bug component, OS, hardware, software version, developer who owns the code, current workload of developers, list of devs actively contributing). Moreover, the author considers launching a Firefox extension to test the app even further through human evaluation as well.
- **What did they learn?** The research brings 3 main contributions: methodology for creating semi-automated bug triaging algorithms, a characterization of performance of different approaches, as well as an actual implementation for such a semi-automated big triaget. The author discovered that with further improvements and tweaking, we can reach an even further bug triager.
- **Relevance to our work:** Another project that we can use to learn about bug triaging; also we can contact maybe the author and receive either an executable or the source code for the application.

Reducing the effort of bug report triage[3]:

- **What were the goals?** The authors had one main goal, and that was to find out if costs related to bug triaging can be reduced if ML algorithms were to be introduced. This hypothesis was backed by 2 research questions, which were if the approach would create recommenders that make accurate recommendations and if the humans can make use of the information provided by the approach.
- **What was the method?** They created a bug triager which can present a list of developers from which someone can choose who is the most suitable to fix that specific issue. Also, they improved a previous research the same authors worked on, thus increasing the precision and recall of the algorithm. They employed machine learning techniques to create the algorithm. Then, in order to test it, they implemented a proxy to the actual web service thats providing the issue repository, thus visualizing results helped conduct the evaluation process.
- **What did they learn?** The authors learned that their approach could actually be used in software projects as its accuracy is reasonable. They also learned that the impact of poor software project management, including bug triaging, on software projects can be quite high.
- **Relevance to our work:** The research presented can help us in understanding better if statistical expetiments are helpful in analyzing tickets, as well as give an insight into how important ticket quality and proper triaging are for any software projet.

Effects of process maturity on quality, cycle time and effort in software product development[14]:

- **What were the goals?** The research tries to investigate the relationship between process maturity, quality, cycle time and effort in software projects. The authors are mainly trying to prove/reject a couple of hypotheses: higher levels of process maturity lead to higher product quality in software projects, higher levels of process maturity are associated with increased cycle time in software products, higher product quality is associated with lower cycle time, higher levels of process maturity lead to increased development efforts in the project and higher product quality is associated with lower development effort.
- **What was the method?** In order to test their hypotheses, the authors examined data coming from 30 projects belonging to the systems integration division from a large IT company. The process improvement data was collected via external divisions and by government agencies to provide independent assessments of the firms software development processes. Then, they analyzed a couple of key variables and see how they interact: process maturity, product quality, cycle time, development effort, product size, domain/data/decision complexity and requirements ambiguity.
- **What did they learn?** They learned that the main features they inspected are additively separable and linear. Moreover, they found out that higher levels of process maturity are associated with significantly higher quality, but also with increased cycle times and development efforts. On the other hand, the reductions in cycle time and effort resulting from improved quality outweigh the marginal increases from achieving more process maturity.
- **Relevance to our work:** It can aid our research through better understanding of how various factors can influence the software development processes, thus giving more value to our own research (i.e. being able to automatically determine issue quality could improve better issue writing guidelines/better rules to be enforced, thus reduced triaging and development costs).

Where should the bugs be fixed?-more accurate information retrieval-based bug localization based on bug reports[38]:

- **What were the goals?** The main goal of the paper was to implement a tool that, based on a bug report, could accurately select a number of files where the developer needs to make the necessary changes to fix the issue. They tried to answer 4 main research questions: how many bugs can be successfully located by BugLocator, does the revised VSM improve bug localization accuracy, does considering similar bugs improve localization accuracy, can bugLocator outperform other similar methods.
- **What was the method?** The tool firstly performed a textual analysis, looking for similarities between the bug report description and the source code files. Then, it analyzed previous bugs in the repository to find the most similar ones, thus being able to find which files ought to be changed. Lastly, it assigned scores to similar files, the ones with bigger sizes obtaining higher scores as they are more likely to contain bugs. They collected the necessary data from Bugzilla projects and then performed the evaluation.

- **What did they learn?** The buglocator can locate a large percentage of bugs analysing just a small set of source code files. Secondly, the revised VSM outperforms the standard VSM. Moreover, similar bugs can improve the localization accuracy only to a certain extent, while the locator outperformed every other competitor on a multitude of projects.
- **Relevance to our work:** We can try to use BugLocator and even include it in our research as it gives a direct correlation between bug reports and development effort (thus reduced costs). Thus, we can infer even more possible hypotheses regarding our issues quality.

Analyzing and Relating Bug Report Data for Feature Tracking[11]:

- **What were the goals?** The main goal of the paper was to analyze the proximity of software features based on modification and problem report data.
- **What was the method?** They employed a method to track features by analyzing and relating bug report data filtered from a release history database. Features are instrumented and tracked, relationships of modification and problem reports to these features are established, and tracked features are visualized to illustrate their otherwise hidden dependencies.
- **What did they learn?** The authors approach suggest first to instrument and track features, then establish the relationships of modification and problem reports to these features and lastly visualize the tracked features for illustrating their non apparent dependencies.
- **Relevance to our work:** This paper can give us insight into how bug reports can influence feature implementation throughout the lifetime of the ticket.

Duplicate Bug Reports Considered Harmful... Really?[8]:

- **What were the goals?** The authors wanted to test two main hypotheses, and these were: duplicate bug reports provide developers with information that was not present in the original report and the information in bug duplicates can improve automated triaging techniques.
- **What was the method?** They collected big amounts of data from the Eclipse open source project in XML form. Then, they ran different kinds of textual and statistical analysis on the data to find answers to their research questions.
- **What did they learn?** They reached the conclusion that bug duplicates contain information that is not present in the master reports. This additional data can be helpful for developers and it can also aid automated triaging techniques (e.g. decide who to assign a bug to).

- **Relevance to our work:** The findings presented in this paper could potentially make us consider duplicate tickets in a different way (i.e. do not treat it as disposable but analyze it and see if it maybe adds more information to the original/master report).

Summarizing Software Artifacts: A Case Study of Bug Reports[29]:

- **What were the goals?** The main goal of the research was to determine if software artifacts could be summarized effectively and automatically so that developers would need only to analyze summaries instead of full software artifacts (i.e. in our case, bug reports).
- **What was the method?** Firstly, they collected data to analyze from 4 different open source projects (Eclipse, Mozilla, Gnome, KDE). Then, they asked the volunteers (i.e. university students) to annotate the bug reports - write a summary of maximum 250 words in their own sentences. These human-produced annotations were then used by algorithms to learn how to effectively summarize a bug report. Afterwards, the authors asked the end users of these bug reports, the software developers, to rate the summaries against the original bug reports.
- **What did they learn?** They learned that existing conversation-based extractive summary generators trained on bug reports produce the best results.
- **Relevance to our work:** As the paper revolves around conversations that are attached to a bug (i.e. comments) instead of the actual description of the bug, we can apply the technique described here maybe to generate summaries and analyze them as well in the overall context of ticket quality (maybe together with some sentiment analysis).

Improving Bug Triage with Bug Tossing Graphs[18]:

- **What were the goals?** The two main goals of the authors, through their bug tossing graph idea (based on the Markov property), were to discover developer networks and team structures, as well as help to better assign developers to bug reports.
- **What was the method?** They analyzed 145.000 bug reports from Eclipse and 300.000 from Mozilla. Then, using statistical analysis on the bug reports in order to find evolution histories and changes throughout the lifetime of the reports, they created the bug tossing graph.
- **What did they learn?** They learned that it takes a long time to assign and toss bugs. Additionally, they learned that their model reduces tossing steps by up to 72% and improved the automatic bug assignment by up to 23%.
- **Relevance to our work:** First paper on importance of tossing and its impact on the software development process. It might prove useful as we shall also inspect tossing and bug assignment in our own research in order to determine ticket quality (i.e. maybe tossing actually reduces the quality of the ticket?).

Towards the Next Generation of Bug Tracking Systems[19]:

- **What were the goals?** The authors wanted to determine what makes for good bug reports and, more specifically, how the reporting environment can be changed to best suit the developers needs (and everyone else involved in the software project).
- **What was the method?** They first launched a survey to developers from Eclipse, Mozilla and Apache open source projects from which they received 175 comments back. Then, they applied a card sort in order to organize the comments into hierarchies to deduce a higher level of abstraction and identify common patterns.
- **What did they learn?** They found 7 main topics of improvement for bug tracking systems: provide tool support for users to collect and prepare information that developers need, find volunteers to translate bug reports filed in foreign languages, provide different user interfaces for each user level and give cues to inexperienced reporters on what to report and best practices, reward reporters when they do a good job, integrate reputation into user profiles to mark experienced reporters, provide powerful and easy to use tools to search bug reports, encourage users to submit additional details (provide tools for merging bugs).
- **Relevance to our work:** This is an important paper to our research as it presents, from open source developers perspective, how bug reporting and bug tracking systems should behave. This research gives us various topics to concentrate on when analyzing the quality of tickets.

Quality of Bug Reports in Eclipse[6]:

- **What were the goals?** The authors were looking to find the most important factors that determine the quality of a bug report.
- **What was the method?** They conducted a survey with Eclipse developers on what they find most important and helpful in a bug report. Another goal was to develop a tool that could automatically measure the quality in bug reports based on the best practices for writing such reports (Eclipse guidelines).
- **What did they learn?** They learned that well written bug reports will most likely get more attention from developers. Moreover, the most important factors which influence ticket quality are stack traces and steps to reproduce. On the other hand, the most harmful information in bug reports are errors in steps to reproduce, incomplete information and wrong observed behavior. Regarding quzilla, they measured that the tool achieved an accuracy of 42% which indicates that it could, on the long term, ensure some standard quality when reporting and creating new tickets.
- **Relevance to our work:** It is mostly the same topic as our research, but we want to use statistical analysis instead of interviews/questionnaires as well as investigate tickets from multiple open source repositories, not just one. However, the findings presented here could serve as clues on where to search for factors that influence ticket quality.

Bug Report Assignee Recommendation using Activity Profiles[25]:

- **What were the goals?** The main goal of the paper is to automatically detect who should be the developers who are most suitable to solve a specific issue.
- **What was the method?** They employed an algorithm using activity profiles (i.e. assign, review, resolve activity of the developer) such that, after detecting the prior experience, developer's role, and involvement in the project, it could recommend who should fix a specific bug. The average accuracy was around 88%, much higher than the LDA-SVM technique.
- **What did they learn?** They found out that their approach was more accurate than previous work done on assignee recommendation, but they think that it might even be improved further by combining their technique with the classic LDA-SVM one.
- **Relevance to our work:** this work complements the research presented in Anvik et al. [2] and it will provide us with valuable insight into how automatic triaging should be performed.

Modeling Bug Report Quality[16]

- **What were the goals?** The authors' goal is to present a model that would determine if a bug report will be triaged in a given amount of time. Moreover, they are trying to propose features that could help composing a good quality bug report.
- **What was the method?** They run an analysis on over 27000 bug reports from the Mozilla Firefox project, selecting a couple of surface features that could be applied:
 - self-reported severity;
 - readability;
 - daily load (i.e. total number of bug reports that need to be dealt with at that point in time);
 - submitter reputation;
 - changes over time.

Afterwards, they ran 4 experiments (i.e. validate the assumptions, find how much post-submission data is necessary, find the optimal resolved by cutoff, evaluate the hypothetical benefit) and analyzed the results to see if their model would be beneficial.

- **What did they learn?** They found out that linear model, even a simple one, can have better-than-chance predictive power. Also, they came to the conclusion that attachment and comment counts are really valuable for triaging the bug faster, compared to, for example, patch count. Moreover, the readability of a bug is important as it proved to influence the duration of the triaging process.

- **Relevance to our work:** the authors' goal was quite similar to ours as they are trying to find out which features make for a good quality bug report. Moreover, they are also using programmatic analysis instead of interviews with developers to produce the results, which is the same approach we are aiming for in our research.

Secret Life of Bugs[4]:

- What were the goals? The paper tries to understand and analyze common bug fixing coordination activities. Another goal of the paper was to analyze the reliability of repositories in terms of software projects coordination and propose different directions on how to implement proper tools.
- What was the method? They executed a field study which was split into two parts:
 - firstly, they did an exploratory case study of bug repos histories;
 - secondly, they conducted a survey with professionals (i.e. testers, developers).

All data and interviews were conducted using Microsoft bug repositories and employees.

- What did they learn? They learned that there are multiple factors which influence the coordination activities that revolve around bug fixing, such as organizational, social and technical knowledge, thus one cannot infer any conclusions only by automatic analysis of the bug repositories. Also, through surveying the professionals, they reached the conclusion that there are 8 main goals which can be used for better tools and practices:
 - probing for ownership;
 - summit;
 - probing for expertise;
 - code review;
 - triaging;
 - rapid-fire emailing;
 - infrequent, direct email;
 - shotgun emails.
- Relevance to our work: this paper is one of the key papers for our research paper.

What makes a good bug report[7]:

- **What were the goals?** The main goal of the paper was to investigate the **quality of bug reports** from a developer's point of view, based on the typical information found in such a report (e.g. stack traces, screenshots).

- **What was the method?** The authors conducted a massive survey with over 450 respondents. The survey was online and it targeted developers from Mozilla, Apache and Eclipse.
- **What did they learn?** The main conclusion of this research paper was that well written bug reports will more likely get the attention of the developers. Thus, including steps to reproduce the bugs or stack traces proved to increase the quality of the bug report. Also, an important achievement reached by the authors was the development of a prototype tool called Cuezilla that could estimate, with an accuracy rate of 3148%, the quality of a bug report.
- **Relevance to our work:** it is relevant to our research in that it provides valuable insight into what makes for a good bug report based on actual professionals' opinions.

Who should fix this bug?[2]:

- **What were the goals?** The authors aimed to create a tool that could automatically assign the bug report to a specific developer based on his/her suitability for that specific task.
- **What was the method?** They applied a supervised machine learning algorithm on the repositories to learn which developers were best suited for specific tasks, thus when a new bug report would come in, a small set of people would be selected. In order to train the algorithm, they looked at Bugzilla repositories and selected the free text form of tickets, trying to label similar ones based on textual similarities. Once the tickets were labeled and grouped for specific developers, the algorithm would then be able to present the triager the set of developers suitable to fix the bug.
- **What did they learn?** The most important lesson learned was that collecting data from bug reports and CVS logs was quite challenging. One of the major reasons why they found this aspect hard was that not all CVS comments referenced the specific bug report id.
- **Relevance to our work** The paper taught us that bug triaging is hard and that there is almost no automated tool that can choose the perfect developers to work on the task. Moreover, the method applied by the authors could prove useful as a learning aid when working with the open source repositories chosen as data sets.

Software Quality The Elusive Target[21]:

- **What were the goals?** The main goal of the paper is to determine what makes for a good quality software project, as well as who are the people in charge of this aspect and how should they approach achieving it.
- **What was the method?** They tried to define quality in software projects and analyze techniques that measure such quality by looking at other models proposed in different other papers (e.g. McCall's quality model, ISO 9126).

- **What did they learn?** They learned that quality is very hard to define and there are various factors which need to be taken into consideration, such as the business model of the company, the type of the software project (e.g. safety critical, financial) or the actors which are involved and how they coordinate the software activities.
- **Relevance to our work:** this paper is a key paper on software quality and it can prove beneficial in our research by giving valuable insights into how software quality can be modeled, thus helping us in selecting good quality open source repositories to work with (i.e. ticket selection and analysis).

Code Quality Analysis in OSS[33]:

- **What were the goals?** The article tries to discuss and examine the quality of the source code delivered by open source projects.
- **What was the method?** They used a set of tools that could automatically inspect various aspects of source code. The authors analyzed the 6th release of the OpenSUSE project and examined only the components, which are defined by C functions in the programs.
- **What did they learn?** The research's results show that Linux applications have high quality code standards that one might expect in an open source repository, but the quality is lower than the one implied by the standard. More than half of the components were in a high state of quality, but on the other hand, most lower quality components cannot be improved only by applying some corrective actions. Thus, even though not all the source code was in an industrial standards shape, there is definitely room for further improvement and open source repositories proved to be of good quality.
- **Relevance to our work:** this work completes the previous paper on software quality in general by looking specifically at quality in open source projects, which will be our main points for data collection.

Analysis of Software Cost Estimation[13]:

- **What were the goals?** The authors are trying to show that poor estimation analysis techniques in software projects will lead to wrong conclusions regarding cost estimation accuracy. Moreover, they also propose a framework for better analysis of software cost estimation error.
- **What was the method?** They approached a real-world company where they conducted analysis on their cost estimation techniques.
- **What did they learn?** They learned that regular, straight-forward types of cost estimation analysis techniques error lead them to wrong conclusions.
- **Relevance to our work:** it showed us that we need to be careful when selecting techniques for cost estimation in our own research.

3.6 Measuring Cost and Waste in Software Projects

Waste Identification[22]:

- **What were the goals?** The paper had two main goals:
 - a means to identify communication waste in agile software projects environments;
 - types of communication waste in agile projects.
- **What was the method?** The authors collaborated with a medium-sized American software company and conducted a series of observations, informal discussions, documents provided by the organization, as well as semi-structured interviews. Moreover, the data collection for waste identification was split into 2 parts:
 - **pre-development:** occurred before the actual implementation begun (e.g. backlog creation);
 - **development:** happened throughout the implementation process (e.g. throughout sprints, retrospectives, sprint reviews, communication media).
- **What did they learn?** They realized the communication waste can be divided into 5 main categories:
 - lack of involvement;
 - lack of shared understanding;
 - outdated information;
 - restricted access to information;
 - scattered information.

Also, they learned that their way of identifying these types of waste was quite efficient and they even recommend it to companies if they'd like to conduct such processes internally.

- **Relevance to our work:** the waste identification process can be applied to our work so that we can identify possible causes to poor quality tickets.

Software Development Waste[32]:

- **What were the goals?** The main goal of the paper was to identify main types of waste in software development projects.
- **What was the method?** They conducted a participant-observation study over a long period of time at Pivotal, a consultancy software development company. They also interviewed multiple engineers and balanced theoretical sampling with analysis to achieve the conclusions.

- **What did they learn?** They found out there are nine main types of waste in software projects:
 - building the wrong feature or product;
 - mismanaging backlog;
 - extraneous cognitive load;
 - rework;
 - ineffective communication;
 - waiting/multitasking;
 - solutions too complex;
 - psychological distress.
 - **Relevance to our work:** this paper complements the previous one on waste identification[22].

Waste in Kanban Projects[17]:

- **What were the goals?** The authors are trying to find the main sources of waste in Kanban software development projects and categorize/rank them based on severity.
- **What was the method?** A controlled case study research was employed in a company called Software Factory. They conducted semi-structured interviews with 5 of the team members both in the beginning in order to collect data as well as at the end of the whole process to categorize the seven types of waste found. Moreover, they also measured the overall success of the project based on Shenar's techniques (first-second-third-fourth; project efficiency-impact on the customer-business success-preparing for the future).
- **What did they learn?** They reached two main findings:
 - they found 7 types of waste throughout the project at various development stages:
 - * partially done work;
 - * extra processes;
 - * extra features;
 - * task switching;
 - * waiting;
 - * motion;
 - * defects.
 - they reached the conclusion that they couldn't explain the success of the project even though waste was found.
- **Relevance to our work:** this work completes the findings from the previous work presented as most of the projects we will work with will be Agile, thus Kanban-based in terms of issue management.

3.7 Sentiment Analysis

Thumbs Up or Thumbs Down[35]:

- **What were the goals?** The main goal of the paper is to detect the overall sentiment transmitted through reviews of various types.
- **What was the method?** The author created an unsupervised machine learning algorithm that was evaluated on more than 400 reviews on Epinions on various kinds of markets (e.g. automobiles, movie). The algorithm implementation was divided into three steps:
 - extract phrases containing adjectives or adverbs;
 - estimate the semantic orientation of the phrases;
 - classify the review as recommended or not recommended based on the semantic orientation calculated at previous step.
- **What did they learn?** One thing the author learned that different categories will yield different results. For example, the automobile section on Epinions ranked much higher, 84%, compared to movie reviews, which had an accuracy of 65.83%. Moreover, most pitfalls of the algorithm could be attributed to multiple factors, such as not using a supervised learning system or limitations of PMI-IR.
- **Relevance to our work:** the method can be applied for extracting sentiments from the tickets (description and comments) we will use in our own research.

Recognizing Contextual Polarity[37]:

- **What were the goals?** The paper's main goal is to find efficient ways to distinguish between contextual and prior polarity.
- **What was the method?** They used a two step method that used machine learning and a variety of features. The first step classified each phrase which had a clue as either neutral or polar, followed by taking all phrases marked in the previous step and giving them a contextual polarity (e.g. positive, negative, both, neutral).
- **What did they learn?** Through the method the authors employed, they managed to automatically identify the contextual polarity. As most papers were only looking at the sentiment extracted from the overall document, they managed to get valuable results from looking at specific words and phrases.
- **Relevance to our work:** when analyzing the description and comments of the ticket, we can use their method for inferring the sentiment transmitted probably more accurately than the technique used in Turney's paper[35].

3.8 Conclusion

4 Proposed Approach

state how you propose to solve the software development problem. Show that your proposed approach is feasible, but identify any risks.

5 Work Plan

show how you plan to organize your work, identifying intermediate deliverables and dates.

References

- [1] John Anvik. Automating bug report assignment. pages 937–940, 2006.
- [2] John Anvik, Lyndon Hiew, and Gail C Murphy. Who should fix this bug? pages 361–370, 2006.
- [3] John Anvik and Gail C Murphy. Reducing the effort of bug report triage: Recommenders for development-oriented decisions. *ACM Transactions on Software Engineering and Methodology (TOSEM)*, 20(3):10, 2011.
- [4] Jorge Aranda and Gina Venolia. The secret life of bugs: Going past the errors and omissions in software repositories. pages 298–308, 2009.
- [5] Adrian Bachmann and Abraham Bernstein. Software process data quality and characteristics: a historical view on open and closed source projects. pages 119–128, 2009.
- [6] Nicolas Bettenburg, Sascha Just, Adrian Schröter, Cathrin Weiß, Rahul Premraj, and Thomas Zimmermann. Quality of bug reports in eclipse. pages 21–25, 2007.
- [7] Nicolas Bettenburg, Sascha Just, Adrian Schröter, Cathrin Weiss, Rahul Premraj, and Thomas Zimmermann. What makes a good bug report? pages 308–318, 2008.
- [8] Nicolas Bettenburg, Rahul Premraj, Thomas Zimmermann, and Sunghun Kim. Duplicate bug reports considered harmful really? pages 337–345, 2008.
- [9] Nicolas Bettenburg, Rahul Premraj, Thomas Zimmermann, and Sunghun Kim. Extracting structural information from bug reports. pages 27–30, 2008.
- [10] Nicolas Bettenburg, Stephen W Thomas, and Ahmed E Hassan. Using fuzzy code search to link code fragments in discussions to source code. pages 319–328, 2012.

- [11] Michael Fischer, Martin Pinzger, and Harald Gall. Analyzing and relating bug report data for feature tracking. 3:90, 2003.
- [12] Narasimhaiah Gorla, Toni M Somers, and Betty Wong. Organizational impact of system quality, information quality, and service quality. *The Journal of Strategic Information Systems*, 19(3):207–228, 2010.
- [13] Stein Grimstad and Magne Jørgensen. A framework for the analysis of software cost estimation accuracy. pages 58–65, 2006.
- [14] Donald E Harter, Mayuram S Krishnan, and Sandra A Slaughter. Effects of process maturity on quality, cycle time, and effort in software product development. *Management Science*, 46(4):451–466, 2000.
- [15] Bernd Heinrich, Marcus Kaiser, and Mathias Klier. Metrics for measuring data quality foundations for an economic data quality management. pages 87–94, 2007.
- [16] Pieter Hooimeijer and Westley Weimer. Modeling bug report quality. pages 34–43, 2007.
- [17] Marko Ikonen, Petri Kettunen, Nilay Oza, and Pekka Abrahamsson. Exploring the sources of waste in kanban software development projects. pages 376–381, 2010.
- [18] Gaeul Jeong, Sunghun Kim, and Thomas Zimmermann. Improving bug triage with bug tossing graphs. pages 111–120, 2009.
- [19] Sascha Just, Rahul Premraj, and Thomas Zimmermann. Towards the next generation of bug tracking systems. pages 82–85, 2008.
- [20] Dongsun Kim, Yida Tao, Sunghun Kim, and Andreas Zeller. Where should we fix this bug? a two-phase recommendation model. *IEEE transactions on software Engineering*, 39(11):1597–1610, 2013.
- [21] Barbara Kitchenham and Shari Lawrence Pfleeger. Software quality: the elusive target [special issues section]. *IEEE software*, 13(1):12–21, 1996.
- [22] Mikko Korkala and Frank Maurer. Waste identification as the means for improving communication in globally distributed agile software development. *The Journal of Systems and Software*, 95:122–140, 2014.
- [23] Yang W Lee and Diane M Strong. Knowing-why about data processes and data quality. *Journal of Management Information Systems*, 20(3):13–39, 2003.
- [24] Yang W Lee, Diane M Strong, Beverly K Kahn, and Richard Y Wang. Aimq: a methodology for information quality assessment. *Information & management*, 40(2):133–146, 2002.
- [25] Hoda Naguib, Nitesh Narayan, Bernd Brügge, and Dina Helal. Bug report assignee recommendation using activity profiles. pages 22–30, 2013.

- [26] R Ryan Nelson, Peter A Todd, and Barbara H Wixom. Antecedents of information and system quality: an empirical examination within the context of data warehousing. *Journal of management information systems*, 21(4):199–235, 2005.
- [27] Leo L Pipino, Yang W Lee, and Richard Y Wang. Data quality assessment. *Communications of the ACM*, 45(4):211–218, 2002.
- [28] Victor R Prybutok, Xiaoni Zhang, and Sherry D Ryan. Evaluating leadership, it quality, and net benefits in an e-government environment. *Information & Management*, 45(3):143–152, 2008.
- [29] Sarah Rastkar, Gail C Murphy, and Gabriel Murray. Summarizing software artifacts: a case study of bug reports. pages 505–514, 2010.
- [30] Thomas C Redman. The impact of poor data quality on the typical enterprise. *Communications of the ACM*, 41(2):79–82, 1998.
- [31] Adrian Schroter, Adrian Schröter, Nicolas Bettenburg, and Rahul Premraj. Do stack traces help developers fix bugs? pages 118–121, 2010.
- [32] Todd Sedano, Paul Ralph, and Cécile Péraire. Software development waste. pages 130–140, 2017.
- [33] Ioannis Stamelos, Lefteris Angelis, Apostolos Oikonomou, and Georgios L Bleris. Code quality analysis in open source software development. *Information Systems Journal*, 12(1):43–60, 2002.
- [34] Diane M Strong, Yang W Lee, and Richard Y Wang. Data quality in context. *Communications of the ACM*, 40(5):103–110, 1997.
- [35] Peter D Turney. Thumbs up or thumbs down?: semantic orientation applied to unsupervised classification of reviews. pages 417–424, 2002.
- [36] Richard Y Wang. A product perspective on total data quality management. *Communications of the ACM*, 41(2):58–65, 1998.
- [37] Theresa Wilson, Janyce Wiebe, and Paul Hoffmann. Recognizing contextual polarity in phrase-level sentiment analysis. pages 347–354, 2005.
- [38] Jian Zhou, Hongyu Zhang, and David Lo. Where should the bugs be fixed?-more accurate information retrieval-based bug localization based on bug reports. pages 14–24, 2012.